FYS-STK3155 - Project 2 Autumn 2019[5]

Johannes A. Barstad, Olve Heitmann

November 10, 2019

Abstract

In this paper we discuss and implement logistic regression (LR), as well as artificial neural networks (ANNs) for use in regression as well as classification problems. The LR and ANN models are trained using batch gradient descent, combined with backpropagation for the ANNs. To test our models we use simulated Franke Function data for regression, and Credit Card data from I-Cheng, Yeh. and Che-hui, Lien's article The comparisons of data mining techniques for the predictive accuracy of probability of default of credit card clients[10] *from 2008. For the simulated Franke Function data our results neural network are able to reproduce results similar to those obtained by OLS and manual feature engineering (polynomials and interaction terms)[1] with a simple two-dimensional input. For the credit card data our results are similar for both LR and ANNs to those obtained by I-Cheng and Che-hui.*

1 Introduction

The main aim of this project was to study classification and regressions problems, including implementing Logistic Regression (LR) as well as implementing and training a Multi Layer Perceptron (MLP) Neural Network using gradient descent and backpropagation. We then apply our models for i) predicting outputs from simulated data using the Franke Function, as well as ii) Predicting credit card defaults.

Dataset i) is our main test subject for the regression task. Our goal here was to study whether or not our MLP were able to reproduce and/ or improve upon the results we obtained in project 1 through manual feature engineering (i.e. the creation of polynomials and interaction terms) simply by taking in a two-dimensional input.

Our study on credit card defaults builds on the work of I-Cheng Yeh and Che-hui Lien[10], which outlines a comparison of various data mining techniques for predicting credit card defaults in the Taiwanese credit card market. I-Cheng and Che-hui's paper focus on a consumer payment dataset from October, 2005, retrieved from an "important bank in Taiwan"[10], as well as reviewing other relevant literature on credit card customers. Our focus here was to test whether or not we could obtain similar results for our implemented LR and MLP models in a classification task.

To justify our various choices made during the implementation we start of by discussing classification, LR and ANNs in a general theory section. For a general discussion on regression techniques, see our first project[1].

2 Theory

This theory section is a summary of what we believe are the most relevant discussions pertaining to our implementation of Logistic Regression and Artificial Neural Networks. The section is based upon lecture notes provided in the course[8,7,9], chapter 4 and 11 in Hastie et. al[6], and chapter 1, 2 and 3 of Michael A. Nielsens book[11].

2.1 On classification problems

Classification problems, as opposed to regression, are concerned with taking on the form of discrete variables (i.e. categories). The perhaps most common, and definitely simplest classification problem, is the *binary classification problem*, where we have two possible outcomes. It is here common to encode the two possible outcomes as

$$y = \begin{bmatrix} 0 & \text{outcome1} \\ 1 & \text{outcome2} \end{bmatrix}$$

In the practical application part of this project we will study the binary case, namely credit card client defaults (two possible outcomes, the client ended up defaulting or the client did not end up defaulting) with a dataset from UCI's Machine

Learning Repository [14].

In general however, the problem can be described as follows: We aim to study the case where the dependent variables (i.e. the responses or outcomes) are discrete and only can take on value from $k = 0, \ldots, K - 1$ — i.e. K classes. We aim to predict the output classes based on *the design matrix* $\hat{\mathbf{X}} \in \mathbb{R}^{n \times p}$. The design matrix consists of n-samples, where each contain data on p features (predictors).

On the use of regressions techniques in classification

One way to employ regressions techniques in classification problems is to have a sign function for mapping the output of the regressor to the discrete values in our list of possible outcomes (e.g. for the binary classification problem to $\{0, 1\}$). An example of such a sign function could be $f(s_i) = sign(s_i) = 1$ if $s_i \geq 0$ and 0 otherwise. This approach is what historically has been called *the perceptron model* in machine learning literature. [8]

Another way to use regression techniques to solve and describe classification problems is to interpret the output of the regression model as a *probability* of a given category. This can be extremely useful in real world applications where we are interested in the expected value of some variable that *depends* on the outcome of the problem we are trying to solve. One example of this could be a banks expected loss on credit debts over a given period, which would be given by *probability of loss* $\times$ *loss* (assuming "binary loss" — i.e. you either default all the money you owe, or nothing at all).
This approach to using regression techniques in classification problems are sometimes referred to as a "soft" classifier — we will study one example of this type of classifier, the *logistic regression*.

2.2 The Sigmoid function

The Sigmoid (logit) function is defined as:

$$f(z) = \frac{1}{1 + \exp{-z}}$$

with the first order derivative given by:

$$f'(z) = f(z)(1 - f(z))$$

Moving on we will refer to the sigmoid function as $\sigma(z)$. Note also that $1 - \sigma(z) = \sigma(-z)$, and that

the function will output a number $(0, 1)$. This is an important attribute making it possible to interpret the output of a sigmoid function as the probability of a binary event.

2.3 The Relu activation function

The Relu activation is a very simple activation function, which has gained popularity due to its good performance in practical application. It takes the form of

$$relu(x) = \max(0, x)$$

The derivative is set to $\begin{cases} 1 \text{ if } x > 0 \\ \text{else } 0 \end{cases}$

2.4 Logistic Regression

For the binary case, logistic regression let us model the probability that a sample is of type "outcome2", as opposed to "outcome1". The binary case model has the form:

$$\log \frac{p(\hat{\beta}\hat{x})}{1 - p(\hat{\beta}\hat{x})} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \ldots \beta_p x_p$$

Logistic regression is widely used in biostatistical applications where binary problems occur frequently (patients survive or not, have heart disease or not, etc.) [6]

2.4.1 Fitting Logistic Regression Models

Logistic Regression models are typically fit by *maximum likelihood*, using the conditional likelihood of G given X, where G can be interpreted as the class of the observed sample, and X the observed explanatory factors.
Maximimum likelihood in general is not discussed in detail here, but can be thought of as methods were we aim at maximizing the probability of seeing the observed data — for a more mathematical rigorous explanation, see e.g. Hastie et. al chapter 8. [6]
As $Pr(G|X)$ completely specifies the conditional distribution, the *multinomial* distribution is appropriate (Hastie. et al page 120 [6]). The log-likelihood is then given by

$$l(\theta) = \sum_{i=1}^{N} \log p_{g_i}(x_i; \theta)$$

where $p_{g_i}(x_i; \theta) = Pr(G = k|X = x_i; \theta)$. Coding $y_i = 1$ when $g_i = 1$, and $y_i = 0$ when $g_i = 2$, and denoting $p_1(x; \theta) = p(x; \theta)$ and $p_2(x; \theta) = 1 - p(x; \theta)$

we can write

$$l(\beta) = \sum_{i=1}^{N} \{y_i \log p(x_i; \beta) +)(1 - y_i) \log (1 - p(x_i; \beta))\}$$
$$= \sum_{i=1}^{N} \{y_i \beta^T x_i - \log (1 + \exp^{\beta^T x_i})\}$$

assuming that the vector of inputs x_i includes the constant term 1 to accommodate the intercept. To maximize the log-likelihood, we set the derivatives to 0. The score equations are then given by

$$\frac{\partial l(\beta)}{\partial \beta} = \sum_{i=1}^{N} x_i(y_i - p(x_i; \beta)) = 0$$

which are $p + 1$ equations *nonlinear in β*

2.5 Neural Networks

Neural networks (NNs) are a class of learning methods developed originally developed separately in two different fields — statistics and artificial intelligence. The two fields however ended up creating what is essentially identical models[6].

NNs work by extracting linear combinations of the inputs as a new set of *derived features*, and then modeling the target as a nonlinear function of the *derived features* (In simpler statistical models the construction of these features is typically more manual, and is often referred to as *feature engineering*). They are constructed as a *two-stage* regression or classification model, with the possible of multiple quantitative final outputs. For K-class classification we typically want K outputs, with the kth output modeling the probability of the sample being in class k. For regression problems 1 output suffice. The two stages can be interpreted as passing the inputs into the hidden layers, and extracting the output layer from the last hidden layer. For training our neural networks we will apply *gradient descent-* and *backpropagation-algorithms* to minimize our chosen cost functions.

2.5.1 Choosing cost functions

This subsection is primarily based on C. Piech and M. Sahamis discussion on Parameter Estimation from May 2017[12].

In general, it is vital that the cost function is differentiable, as it is going to be used in the context of gradient descent. Further on, it is useful for these derivative terms to be simplistic. We chose the following cost functions as they satisfy these conditions clearly, as well as they have a good theoretical backing from maximum likelihood estimation of parameters.

For the regression part we apply *quadratic loss*;

$$C = \frac{1}{2} \sum_{i}^{n} (\hat{y}_i - y_i)^2$$

and for the classification problem we apply *logistic loss*

$$C = -\frac{1}{n} \sum_{x} [y \ln a + (1 - y) \ln(1 - a)]$$

The quadratic loss is a natural choice of cost function when handling regression problems where the range of the target variable is potentially $(-\inf, \inf)$. It stems from the maximum likelihood estimation based on the normal distribution, when samples are assumed IID. Intuitively it also seems like a natural choice as it is putting a larger penalty to larger deviations and is symmetric in its penalty, suggesting that both positive and negative deviations are equally important. Which is suitable in this situation.

The logistic loss is a natural choice when our targets range takes on a value in $(0, 1)$. It stems from Bernoulli maximum likelihood estimation. Since the $\ln(a)$ is not defined when $a = 0$, in practice one often ads an additional small term to a in order to prevent this effect from impacting the general implementation. We can also easily see that inserting a small value of a (close to 0) while $y = 1$ will return a large value due to $\ln(a)$ will be large and negative. And equally when $y = 0$, and a is close to 1, $\ln(1 - a)$ will be large and negative. While when they are close, the value is small. Thus the penalty grows exponentially as the deviation from the truth grows, which is desirable for a cost function. To combat the negative values from the logarithmic terms, we also add a minus sign in front, as we are minimizing this expression.

2.5.2 Gradient descent and minimizing cost functions

The fundamental idea of the gradient descent is that a function $F(\vec{x})$ decreases fastest if one moves $\vec{x}$ in the direction of the negative gradient:

$$\partial f(X) = \nabla_X f(X) \partial X$$

where ∇_X is the gradient of f with respect to X, i.e.

$$\nabla_X \equiv \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \cdots, \frac{\partial f}{\partial x_m} \right)$$

This idea stems from a well known linear algebraic relation, the inner product between two vectors of fixed lengths:

$$\vec{w}^T \vec{v} = |w||v| \cos \theta$$

This product can be shown to be at its maximum when the two vectors are aligned (i.e. when $\theta = 0$), and consequently at its minimum when the two vectors point in opposite directions.

In summary, finding the minimum of our scalar function requires us to find a turning point — i.e. a point where the gradient will be 0. As most of the cost functions desirable to apply for neural networks doesn't have closed form solutions, it is often not possible to simply solve $\nabla_X f(X) = 0$.

Iterative solutions
When closed form solutions are unavailable we begin with an initial "guess" for the optimal X and refine it iteratively until the correct value is obtained — One should here however note the "risk" of discovering mere local minima. For *squared error loss*

$$Loss = \frac{1}{T} \sum_i div(Y_i, d_i)$$

Minimize wrt $\{w_{ij}^l, b_j^l\}$ intialize all weights and biases $\{w_{i,j}^k\}$ For every layer k, for all i, j update $w_{i,j}^k = w_{i,j}^k - \eta \frac{\partial Loss}{\partial w_{i,j}^k}$ until loss has converged

One starts of with an initial guess LOL for the minimum value of F, and compute new approximations according to

- Initialize x^0, $k = 0$

- While $||\nabla_x f(x^k)|| > \epsilon$:

$$x^{k+1} = x^k - \eta^k \nabla_x f(x^k)^T$$

where η^k is the "step size" (learning rate), and ϵ is some threshold "Taking a step against the derivative"

2.5.3 Implementing backpropagation

Backpropagation was first introduced in a famous 1986 paper by David Rumelhart, Geoffrey Hinton, and Ronald Williams. The paper describes several neural networks where backpropagation works far faster than earlier approaches, making it possible to use neural nets to solve problems which had previously been insoluble. Backpropagation is now the main workhorse of learning in neural networks.

Assumptions for employing backpropagation

To successfully apply backpropagation we need to make two assumptions regarding the cost function:

1. **The cost function can be written as an average of the cost functions of the individual training samples:**

$$C = \frac{1}{n} \sum_i^n C_i$$

Backpropagation let us compute the partial derivatives $\frac{\partial C_i}{\partial w}$ and $\frac{\partial C_i}{\partial b}$ for a single training sample. We then need to "recover" $\frac{\partial C}{\partial w}$ and $\frac{\partial C}{\partial b}$ by averaging over the training samples — if this is not possible due to the cost function we choose, backpropagation will not give good approximations to the optimum parameter values.

2. **Cost can be written as a function of the outputs from the neural network** As the parameters we are able to affect (weights, biases) only have a direct effect on the outputs from the neural network, the cost function needs to be able to be written as a function of these outputs.

The fundamental equations of backpropagation
Backpropagation is based around four fundamental equations:

$$\delta_j^L = \frac{\partial C}{\partial a_j^L} \sigma'(z_j^L). \qquad \text{(BP1)}$$

$$\delta^L = (a^L - y) \odot \sigma'(z^L). \qquad \text{(BP2)}$$

$$\frac{\partial C}{\partial b_j^l} = \delta_j^l. \qquad \text{(BP3)}$$

$$\frac{\partial C}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l. \qquad \text{(BP4)}$$

An equation for the error in the output layer

$$\delta_j^L = \frac{\partial C}{\partial a_j^L} \sigma'(z_j^L)$$

1. **Input x:** Set the corresponding activation a^l for the input layer

2. **Feed forward:** For each $l = 2, 3, \ldots, L$ compute $z^l = w^l a^{l-1} + b^l$, and $a^l = \sigma(z^l)$

3. **Output error σ^L:** Compute the vector $\sigma^L = \nabla_a C \odot \sigma'(z^L)$

4. **Backpropagate the error:** For each $l = L - 1, L - 2, \ldots, 2$ compute $\delta^l = ((w^{l+1})^T \delta^{l+1}) \odot \sigma'(z^l)$

5. **Output:** The gradient of the cost function is given by $\frac{\partial C}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l$ and $\frac{\partial C}{\partial b_j^l} = \delta_j^l$

We compute the error vectors δ^l backward, starting from the final layer — hence *backpropagation.* the backward movement is a consequence of the fact that the cost is a function of outputs from the network. To understand how the cost varies with earlier weights and biases we need to repeatedly apply the chain rule, working backward through the layers to obtain usable expressions. In practice, it's common to combine backpropagation with a learning algorithm such as stochastic gradient descent, in which we compute the gradient for many training examples. In particular, given a mini-batch of m training examples, the following algorithm applies a gradient descent learning step based on that mini-batch

1. Input the set of training samples

2. Repeat step 1-4 of the back propagation algorithm for each training sample x, that is:

   (a) **Input x:** Set the corresponding input activation $a^{(x,l)}$

   (b) **Feed forward:** For each $l = 2, 3, \ldots, L$ compute $z^{(x,l)} = w^l a^{(x,l-1)} + b^l$, and $a^{(x,l)} = \sigma(z^{(x,l)})$

   (c) **Output error $\sigma^{(x,L)}$:** Compute the vector $\sigma^{(x,L)} = \nabla_a C_x \odot \sigma'(z^{(x,L)})$

   (d) **Backpropagate the error:** For each $l = L - 1, L - 2, \ldots, 2$ compute $\delta^{(x,l)} = ((w^{l+1})^T \delta^{(x,l+1)}) \odot \sigma'(z^{(x,l)})$

3. **Gradient descent:** For each $l = L, L - 1, \ldots, 2$ update the weights and biases according to the following rules:

   - $w^l \rightarrow w^l - \frac{n}{m} \sum_x^n \delta^{(x,l-1)} (a^{(x,l-1)})^T$

   - $b^l \rightarrow b^l - \frac{n}{m} \sum_x^n \delta^{(x,l-1)}$

Example: NN Loss development during training, classification

Loss development during training
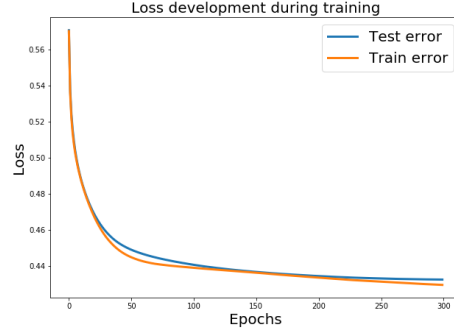
To implement the full stochastic gradient descent algorithm we also create an outer loop to generate mini-batches of training samples, as well as an outer loop stepping through the multiple epochs of training

2.5.4 Regularization in Neural Networks

To prevent overfitting and stabilizing the network we apply *regularization* (for a more in-depth discussion of regularization, see our first project, project 1[1]). This is done by adding a regularization term $\frac{\lambda}{2n} \sum_w w^2$ to the cost function C_0:

$$C = C_0 + \frac{\lambda}{2n} \sum_w w^2 \quad \text{(Regularized cost function)}$$

The effect of regularization, intuitively, is to make the network "prefer" to learn small weights, all else equal. Taking the partial derivatives of the regularized cost function above, we get:

$$\frac{\partial C}{\partial w} = \frac{\partial C}{\partial w} + \frac{\lambda}{n} w$$
$$\frac{\partial C}{\partial b} = \frac{\partial C_0}{\partial b}$$

3 Datasets

3.1 Data generated by the Franke Function

For the regression part of the implementation we again use the data generated by the Franke Function from Project 1[1]. We here chose to use the non-noisy version, as one of our goals was to study whether or not the ANN actually was able to produce feature engineering similar to what we implemented in Project 1 using polynomials and interaction terms.

3.2 Default of credit card clients

In the classification part of the project we will study a cross sectional dataset containing credit

card clients who either have, or have not defaulted. The time horizon is one month from the last observation of the proposed explanatory variable, until the observation of the dependent variable. Note however that some of the explanatory variables are observed even before one month before the observation of the dependent variable — the dataset could thus be constructed as a *time series dataset*. We did not however do the data wrangling necessary to model it that way here. For a more detailed overview of the dataset, see table 4 in the appendix.

4 Application and results

4.1 Scoring metrics for classification problems

To properly assess the results from applying our models to the datasets discussed in section 3 we use the following metrics, as well as R^2 for the regression problem. For a discussion on R^2, see our report for Project 1[1].

Accuracy score
The accuracy score for binary classification problems are given by

$$Accuracy = \frac{\sum_{i=1}^{n} I(t_i = y_i)}{n}$$

where t_i is the guessed target, y_i I is the indicator function, equaling 1 if $t_i = y_i$ and equaling 0 otherwise. As we can see, a perfect classifier will have an accuracy score of 1, while a classifier that predicts every target wrong will have an accuracy score of 0. It also follows that a classifier that is correct exactly half of the time will have an accuracy score of 0.5, making interpretation easy.

Area ratio
Because of the large class imbalance in the credit card data set, where 87.88% of samples are non-risky, we cannot rely on accuracy alone as a measure of model performance, as it may seem better than it actually is. Considering that a completely generic algorithm that is predicting only 0 will give a an accuracy score of 87.88%, which is many cases would be considered impressive.

From a business perspective, one would also want to reduce the number of false positive cases, which will classify potent customers as risky, and thus reduce income for the business. Thus, an alternative evaluation metric is proposed. The Area Ratio is a metric to evaluate models based on three curves. The curves are formed by ranking the customers based on the probability of being a risky customer. The baseline curve suggest random behavior. The model curve uses the scores from the model to rank the customers, and accumulates risky customers as they occur. The optimal curve is a curve where all the risky customers appear before the non-risky customers. We then compare the optimal curve vs the model curve, using the baseline curve as the baseline.

The final metric is given by computing two areas and calculating the ratio between them. The first ratio is the area between the baseline and the optimal curve. The second is the area between the model curve and the baseline curve. An optimal score is 1, where the model curve equals the optimal curve. In this case, the false positives are minimal.

4.2 Results and comparison

4.2.1 Regression on the Franke Function data

ANN Hyperparameter search

General remark on hyperparameter optimization for neural networks: Since neural network models tends to have a large set of hyperparameters, from number of neurons, regularization, learning rate etc, and take a fairly long time to train, hyperparameter tuning can take a long time. Possible remedies to this is to use multiple computers in a cloud environment and train many networks in parallel across different configurations. With domain expertise, there is also possible to limit the search space.

We use a grid search approach to determine the hyperparameter settings.

For the Franke data, we see that a learning rate of 0.1 to 0.01 performs best. There is little difference in the error for changes in lambda, but the best result is achieved at λ 0.1.

Table 1: ANN Results compared to Project 1 OLS[1]

	Test	Train
Barstad, Heitmann, ANN (Project 2)	0.999	0.999
Barstad, Heitmann, OLS (Project 1)	0.998	0.998

4.2.2 Classification on the Credit Card data

We see that the ANN performes well on the credit data, in line with the results from I-Cheng and Che-hui's paper[10]. The lift is substantial compared to the baseline, as seen from the lift charts. We see good performance in both test and train, suggesting a good generalization to the underlying data generating process. The good lift ensures that the credit company can reduce their risk at a relatively low cost, by using the model to filter out the risky customers, while only filtering out a small number of potent customers in the process.

For the case of logistic regression, we see a reduction in performance, and thus a higher cost in terms of lost potent customers related to the process of filtering out customers with high risk. We use the same approach for model selection as with the Franke data to select the hyperparameters for the models.

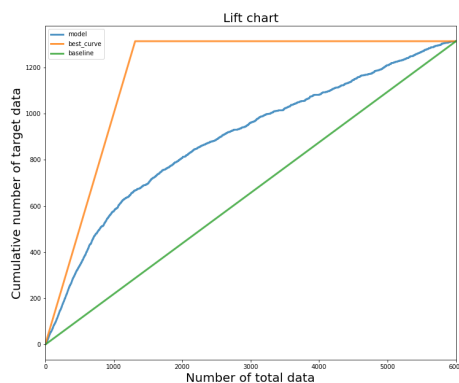**LR Hyperparameter search**

ANN Lift Chart for Test-data

ANN Hyperparameter search

LR Lift Chart for Test-data

Table 2: LR Results compared to I-Cheng, Che-hui [10]

	Area ratio	Accuracy
I-Cheng, Che-hui, Test	0.44	0.82
I-Cheng, Che-hui, Train	0.41	0.80
Barstad, Heitmann, Test	0.46	0.81
Barstad, Heitmann, Train	0.45	0.81

Table 3: ANN Results compared to I-Cheng, Che-hui [10]

	Area ratio	Accuracy
I-Cheng, Che-hui, Test	0.54	0.83
I-Cheng, Che-hui, Train	0.55	0.81
Barstad, Heitmann, Test	0.55	0.82
Barstad, Heitmann, Train	0.58	0.82

5 Conclusion and final remarks

5.1 Conclusion: Assessing the different models

OLS vs ANNs for regression
In terms of accuracy (estimated with test R^2) our ANN implementation performed similar to our OLS implementation from project 1 [1]. The main differences can thus be summarized in that

- Our OLS-implementation employed "manual" feature engineering, which would have been very time consuming had we not known the exact function generating the data

- The OLS-implementation is far easier to interpret

Which model is "best" then depends on the use case. In most scenarios we won't know the exact function that generates the data (if we indeed did — modeling it would be nonsensical), and feature engineering would be very time consuming, and perhaps not even possible with a similar level of success that we had with the ANN. On the other hand, OLS leaves us with an easier to interpret model (albeit that is reduced somewhat by the large number of regressors needed to reach a good fit).

Logistic Regressions vs ANNs for classification

A logistic regression model specifies that an appropriate function of the estimated probability of an event is a linear function of the observed values of the available explanatory variables [10]. The model can produce a simple and fairly interpretable formula for classification, however it is typically harder to properly model non-linear and interactive effects between the observed explanatory variables.

Artificial neural networks can use non-linear equations to develop relationships between input and output variables through what is called a *learning process*. In our study we, as well as I-Cheng and Che-hui [10], apply a back propagation feed-forward neural network. As we have well labeled data (i.e. customers who defaulted, and customers who did not default), we were able to use supervised learning. ANNs can handle non-linear and interactive effects of explanatory variables, but does not have an easily interpreted formula for classification.

5.2 Final remarks

This paper examined various regression and classification techniques, applying Logistic Regression and Artificial Neural Networks on the datasets generated by the Franke-Function, and the Credit Card data provided in I-Cheng and Che-hui's paper [10].

5.3 Feedback for future versions of this course

When starting this project, we expected to be able to use the feedback from Project 1 to improve upon our delivery from the last project. As we still haven't gotten any feedback, we were not able to do so. This is in our opinion very unfortunate. Additionally, we are under the impression that some students have received feedback from the first project in devilry, and thus have been able to take into account your remarks to improve their report for Project 2. If this is the case, scoring and feedback on this project (i.e. project 2) should reflect this difference in possibilities to improve upon ones report.

6 Appendix

6.1 Section 2 — Network Architecture

For the credit card data, we use the following network structure. We have 2 hidden layers of size 100 neurons. We use relu activation functions. The output layer has a single neuron and sigmoid activation. We use a large batch size of 512, and train for 300 epochs. We add a small lambda of 0.001 and use a learning rate of 0.01.

For the Franke data, we use the following network structure. We have 5 hidden layers of size 300, 100, 50, 20 and 10 neurons respectively. We use relu activation functions on these layers. The output layer has a single neuron and linear activation. We use a large batch size of 512, and train for 1000 epochs. We use no regularization. Learning rate is set to 0.01.

6.2 Section 3 — Datasets

Table 4: Default of credit card clients attribute information

Explanatory variable	Explanation
X_1	Amount of the given credit (NT dollar): Individual consumer credit and family credit.
X_2	Gender (1 = male; 2 = female).
X_3	Education (1 = graduate school; 2 = university; 3 = high school; 4 = others).
X_4	Marital status (1 = married; 2 = single; 3 = others).
X_5	Age (year).
X_6 - X_{11}	(Monthly) Past payment status. Scale: -1 = pay duly; 1 = delay 1 month; 2 = delay 2 months
X_{12} - X_{17}	(Monthly) Amount of bill statement (NT dollar).
X_{18} - X_{23}	(Montly) Amount of previous payment (NT dollar).

## 6.3   Section 4 — Results (Misc plots)

LR Lift Chart for Train-data, classification


ANN Lift Chart for Train-data, classification

LR Loss development during training, classification

**LR Regularization effect on area ratio, classification**

ANN Regularization effect on area ratio, classification

LR Regularization effect on weight norm, classification

ANN Regularization effect on weight norm, classification

Regularization effect on L2 weight norm

[Figure: Plot titled "Regularization effect on L2 weight norm", with x-axis "Lambda" (0.001, 0.01, 0.1, 1, 10) and y-axis "Weight L2 norm" ranging 0 to 140.]

References

[1] Johannes A. Barstad and Olve Heitmann. Project1 github repository. https://github.com/Barstad/fys_stk3155/tree/master/PROJECT1.

[2] Johannes A. Barstad and Olve Heitmann. Project2 github repository. https://github.com/Barstad/fys_stk3155/tree/master/PROJECT2.

[3] Derek Bingham. Virtual library of simulation experiments: Test functions and datasets. franke's function. https://www.sfu.ca/~ssurjano/franke2d.html.

[4] M. Chou. Cash and credit card crisis in taiwan, 2006.

[5] Norway Department of Physics, University of Oslo. Project 2 assignment description. https://compphysics.github.io/MachineLearning/doc/Projects/2019/Project2/pdf/Project2.pdf.

[6] Trevor Hastie, Robert Tibshirani, and Jerome H. Friedman. *The Elements of Statistical Learning*. Springer, 2008.

[7] Morten Hjorth-Jensen. Data analysis and machine learning lectures: Optimization and gradient methods. https://compphysics.github.io/MachineLearning/doc/pub/Splines/html/._Splines-bs000.html.

[8] Morten Hjorth-Jensen. Data analysis and machine learning: Logistic regression. https://compphysics.github.io/MachineLearning/doc/pub/LogReg/html/LogReg.html.

[9] Morten Hjorth-Jensen. Data analysis and machine learning: Neural networks, from the simple perceptron to deep learning. https://compphysics.github.io/MachineLearning/doc/pub/NeuralNet/html/NeuralNet-bs.html.

[10] Che-hui Lien I-Cheng, Yeh. The comparisons of data mining techniques for the predictive accuracy of probability of default of credit card clients. https://bradzzz.gitbooks.io/ga-seattle-dsi/content/dsi/dsi_05_classification_databases/2.1-lesson/assets/datasets/DefaultCreditCardClients_yeh_2009.pdf.

[11] Michael A. Nielsen. *Neural Networks and Deep Learning*. Determination Press, 2015.

[12] C. Piech and M. Sahami. 11. parameter estimation, 2017.

[13] USGS. Earthexplorer. https://earthexplorer.usgs.gov/.

[14] USGS. Uci machine learning repository - default of credit card clients data set. https://earthexplorer.usgs.gov/.